

# *Data Wrangling & RMarkdown*

Mark Andrews

Psychology Department, Nottingham Trent University

✉ `mark.andrews@ntu.ac.uk`

# *What is data wrangling*

- ▶ Also known as *data munging*, *data cleaning*, *data pre-processing*, *data preparation*, and so on, *data wrangling* is the process of taking data in its unstructured, messy, or complicated original form and converting it into a clean and tidy format that allows data exploration, visualization, and eventually statistical modelling and analysis to proceed efficiently and relatively effortlessly.
- ▶ Data science's 80/20 rule is that we spend up to 80% of time doing data wrangling, and the rest doing analysis per se.

## *Data frames, tibbles, and reading in data*

- ▶ In everything we do here, and almost always in R generally, data is assumed to be stored in a *data frame*.
- ▶ Data frames are heterogeneous rectangular data structures. Columns are homogeneous vectors, but different columns may be of different types.
- ▶ The tidyverse's style of data frame is known as a *tibble*. This is a regular R data frame but with some superficial and style changes.

## *Reading in data*

- ▶ In practice, it is common to have data in a roughly rectangular format, i.e. with rows and columns, either in text files such as `.csv`, `.tsv`, `.txt`, files.
- ▶ The `readr` package, which is loaded when we load `tidyverse`, allows us to efficiently import data that are in these files.
- ▶ It includes commands
  - ▶ `read_csv` for files where the values on each line are separated by commas
  - ▶ `read_tsv` for files where the values are separated by tabs
  - ▶ `read_delim` for files where the values are separated by arbitrary delimiters such as `'|'`, `':'`, `';`, etc.
  - ▶ `read_table` for files where the values are separated by one or more, and possibly inconsistently many, whitespaces.
- ▶ The `readxl::read_excel` can be used for Excel (`.xlsx`) files.

## Width

```
getOption('width')
```

```
## [1] 80
```

## *Our main example data-set*

```
library(tidyverse)
blp_df <- read_csv("data/blr-trials-short.txt")
blp_df
```

```
## # A tibble: 1,000 x 7
```

```
##   participant lex    spell    resp    rt prev.rt rt.raw
##           <dbl> <chr> <chr>    <chr> <dbl>    <dbl> <dbl>
## 1             20 N      staud    N      977      511    977
## 2              9 N    dinbuss    N      565      765    565
## 3             47 N    snilling    N      562      496    562
## 4            103 N    gancens    N      572      656    572
## 5             45 W     filled    W      659      981    659
## 6             73 W    journals    W      538     1505    538
## 7             24 W     apache    W      626      546    626
## 8             11 W     flake    W      566      717    566
## 9             32 W    reliefs    W      922     1471    922
## 10            96 N     sarves    N      555      806    555
## # ... with 990 more rows
```

## *Glimpsing at data with `glimpse`*

- We can use the `dplyr` command `glimpse` to look at resulting data frame.

```
glimpse(bl_p_df)
```

```
## Rows: 1,000
## Columns: 7
## $ participant <dbl> 20, 9, 47, 103, 45, 73, 24, 11, 32,...
## $ lex         <chr> "N", "N", "N", "N", "W", "W", "W", ...
## $ spell       <chr> "staud", "dinbuss", "snilling", "ga...
## $ resp        <chr> "N", "N", "N", "N", "W", "W", "W", ...
## $ rt          <dbl> 977, 565, 562, 572, 659, 538, 626, ...
## $ prev.rt     <dbl> 511, 765, 496, 656, 981, 1505, 546,...
## $ rt.raw      <dbl> 977, 565, 562, 572, 659, 538, 626, ...
```

- As we can see, there are 1000 rows and 7 variables.

## *The dplyr verbs*

- ▶ The dplyr package provides a set of *verbs* for data wrangling:
  1. `select` for selecting variables
  2. `rename` for renaming variables
  3. `slice` for selecting rows by their indices
  4. `filter` for selecting rows by some condition
  5. `mutate` for adding or modifying variables
  6. `arrange` for sorting variables
- ▶ In addition, there is `summarize` (with `group_by`) for summarizing variables.
- ▶ Each verb performs a major, if focused, task, and they can be chained together using *pipes*.



## *Selecting variables*

- ▶ In our `blp_df` data frames we have 7 variables.
- ▶ Let's say, as is often the case when processing raw data, that we only need some of these.
- ▶ The `dplyr` command `select` allows us to select those we want.

## Selecting variables: Example 1

- Here, we select some variables by name.

```
select(blp_df, participant, lex, resp, rt)
```

```
## # A tibble: 1,000 x 4
##   participant lex    resp    rt
##         <dbl> <chr> <chr> <dbl>
## 1          20 N      N      977
## 2           9 N      N      565
## 3          47 N      N      562
## 4         103 N      N      572
## 5          45 W      W      659
## 6          73 W      W      538
## 7          24 W      W      626
## 8          11 W      W      566
## 9          32 W      W      922
## 10         96 N      N      555
## # ... with 990 more rows
```

## Selecting variables: Example 2

We can select a range of variables by specifying the first and last variables in the range with a `:` between them as follows.

```
select(blp_df, spell:prev.rt)
```

```
## # A tibble: 1,000 x 4
##   spell      resp      rt prev.rt
##   <chr>    <chr> <dbl>   <dbl>
## 1 staud      N      977     511
## 2 dinbuss    N      565     765
## 3 snilling   N      562     496
## 4 gancens    N      572     656
## 5 filled     W      659     981
## 6 journals   W      538    1505
## 7 apache     W      626     546
## 8 flake      W      566     717
## 9 reliefs    W      922    1471
## 10 sarves    N      555     806
## # ... with 990 more rows
```

## Selecting variables: Example 3

We can select a range of variables using indices as in the following example.

```
select(blp_df, 2:5) # columns 2 to 5
```

```
## # A tibble: 1,000 x 4
##   lex    spell    resp    rt
##   <chr> <chr>    <chr> <dbl>
## 1 N      staud      N      977
## 2 N      dinbuss    N      565
## 3 N      snilling  N      562
## 4 N      gancens    N      572
## 5 W      filled     W      659
## 6 W      journals   W      538
## 7 W      apache     W      626
## 8 W      flake      W      566
## 9 W      reliefs    W      922
## 10 N     sarves     N      555
## # ... with 990 more rows
```

## Selecting variables: Example 4

We can select variables according to the character or characters that they begin with.

```
select(blp_df, starts_with('p'))
```

```
## # A tibble: 1,000 x 2
##   participant prev.rt
##   <dbl>      <dbl>
## 1         20      511
## 2          9      765
## 3         47      496
## 4        103      656
## 5         45      981
## 6         73     1505
## 7         24      546
## 8         11      717
## 9         32     1471
## 10        96      806
## # ... with 990 more rows
```

## Selecting variables: Example 5

We can select variables by the characters they end with.

```
select(blp_df, ends_with('t'))
```

```
## # A tibble: 1,000 x 3
##   participant      rt prev.rt
##       <dbl> <dbl>   <dbl>
## 1          20   977     511
## 2           9   565     765
## 3          47   562     496
## 4         103   572     656
## 5          45   659     981
## 6          73   538    1505
## 7          24   626     546
## 8          11   566     717
## 9          32   922    1471
## 10         96   555     806
## # ... with 990 more rows
```

## Selecting variables: Example 6

We can select variables that contain a certain set of characters in any position.

```
select(blp_df, contains('rt'))
```

```
## # A tibble: 1,000 x 4
##   participant    rt prev.rt rt.raw
##   <dbl> <dbl>   <dbl> <dbl>
## 1         20   977     511   977
## 2          9   565     765   565
## 3         47   562     496   562
## 4        103   572     656   572
## 5         45   659     981   659
## 6         73   538    1505   538
## 7         24   626     546   626
## 8         11   566     717   566
## 9         32   922    1471   922
## 10        96   555     806   555
## # ... with 990 more rows
```

## Selecting variables: Example 7

Using matches, the regular expression `^rt|rt$` will match the `rt` if it begins or ends a string.

```
select(blp_df, matches('^rt|rt$'))
```

```
## # A tibble: 1,000 x 3
##       rt prev.rt rt.raw
##   <dbl>   <dbl> <dbl>
## 1    977     511    977
## 2    565     765    565
## 3    562     496    562
## 4    572     656    572
## 5    659     981    659
## 6    538    1505    538
## 7    626     546    626
## 8    566     717    566
## 9    922    1471    922
## 10   555     806    555
## # ... with 990 more rows
```



## Selecting variables: Example 8

To remove a variable, we precede its name with a minus sign.

```
select(blp_df, -participant) # remove `participant`
```

```
## # A tibble: 1,000 x 6
##   lex    spell    resp    rt prev.rt rt.raw
##   <chr> <chr>    <chr> <dbl>   <dbl>   <dbl>
## 1 N      staud      N      977     511     977
## 2 N      dinbuss    N      565     765     565
## 3 N      snilling   N      562     496     562
## 4 N      gancens    N      572     656     572
## 5 W      filled     W      659     981     659
## 6 W      journals   W      538    1505     538
## 7 W      apache     W      626     546     626
## 8 W      flake      W      566     717     566
## 9 W      reliefs    W      922    1471     922
## 10 N     sarves     N      555     806     555
## # ... with 990 more rows
```

## Selecting variables: Example 9

```
select(blp_df, -(2:6))
```

```
## # A tibble: 1,000 x 2
##   participant rt.raw
##   <dbl> <dbl>
## 1      20    977
## 2       9    565
## 3      47    562
## 4     103    572
## 5      45    659
## 6      73    538
## 7      24    626
## 8      11    566
## 9      32    922
## 10     96    555
## # ... with 990 more rows
```

## Selecting variables: Example 9

```
select(blp_df, -contains('rt'))
```

```
## # A tibble: 1,000 x 3
##   lex    spell    resp
##   <chr> <chr>    <chr>
## 1 N      staud      N
## 2 N      dinbuss    N
## 3 N      snilling   N
## 4 N      gancens    N
## 5 W      filled     W
## 6 W      journals   W
## 7 W      apache     W
## 8 W      flake      W
## 9 W      reliefs    W
## 10 N     sarves     N
## # ... with 990 more rows
```

## Selecting variables: Example 10

We can reorder variables using `everything()` as follows.

```
select(blp_df, spell, participant, resp, everything())
```

```
## # A tibble: 1,000 x 7
##   spell      participant resp  lex      rt prev.rt rt.raw
##   <chr>          <dbl> <chr> <chr> <dbl>    <dbl> <dbl>
## 1 staud             20 N      N     977      511    977
## 2 dinbuss           9 N      N     565      765    565
## 3 snilling          47 N      N     562      496    562
## 4 gancens          103 N      N     572      656    572
## 5 filled           45 W      W     659      981    659
## 6 journals          73 W      W     538     1505    538
## 7 apache           24 W      W     626      546    626
## 8 flake            11 W      W     566      717    566
## 9 reliefs          32 W      W     922     1471    922
## 10 sarves           96 N      N     555      806    555
## # ... with 990 more rows
```

## *The `_if`, `_at`, `_all` verb variants*

- ▶ Almost all dplyr verbs have derivatives — either `_if` or `_at` or `_all`, or all of them — that extend their power.
- ▶ For example, `select_if` function allows us to select variables according to properties of their values.

## *Selecting variables: Example 11*

Here, we select all variables that are character vectors.

```
select_if(bl_p_df, is.character)
```

## Selecting variables: Example 12

The following function will return TRUE if the variable is a numeric variable with a mean that is less than 700.

```
has_low_mean <- function(x){  
  is.numeric(x) && (mean(x, na.rm = T) < 700)  
}
```

## Selecting variables: Example 12

Now, we can select variables that meet this criterion as follows.

```
select_if(blp_df, has_low_mean)
```

```
## # A tibble: 1,000 x 3
##   participant      rt prev.rt
##       <dbl> <dbl>   <dbl>
## 1          20   977     511
## 2           9   565     765
## 3          47   562     496
## 4         103   572     656
## 5          45   659     981
## 6          73   538    1505
## 7          24   626     546
## 8          11   566     717
## 9          32   922    1471
## 10         96   555     806
## # ... with 990 more rows
```



## *Renaming variables*

When we select individual variables with `select`, we can rename them too.

```
select(blp_df, subject=participant, reaction_time=rt)
```

```
## # A tibble: 1,000 x 2
##   subject reaction_time
##   <dbl>         <dbl>
## 1         20          977
## 2          9          565
## 3         47          562
## 4        103          572
## 5         45          659
## 6         73          538
## 7         24          626
## 8         11          566
## 9         32          922
## 10        96          555
## # ... with 990 more rows
```

## *Renaming variables*

If we want to rename some variables, and get a data frame with all variables, including the renamed ones, we should use `rename`.

```
rename(blp_df, subject=participant, reaction_time=rt)
```

```
## # A tibble: 1,000 x 7
```

	subject	lex	spell	resp	reaction_time	prev.rt	rt.raw
	<dbl>	<chr>	<chr>	<chr>	<dbl>	<dbl>	<dbl>
## 1	20	N	staud	N	977	511	977
## 2	9	N	dinbuss	N	565	765	565
## 3	47	N	snilling	N	562	496	562
## 4	103	N	gancens	N	572	656	572
## 5	45	W	filled	W	659	981	659
## 6	73	W	journals	W	538	1505	538
## 7	24	W	apache	W	626	546	626
## 8	11	W	flake	W	566	717	566
## 9	32	W	reliefs	W	922	1471	922
## 10	96	N	sarves	N	555	806	555

```
## # ... with 990 more rows
```

## *Renaming all with rename\_all*

The `rename_all` function allows us to rename all the variables using some renaming function, e.g.,

```
rename_all(bl_df, ~str_replace_all(., '\\\\.', '_'))
```

```
## # A tibble: 1,000 x 7
```

##	participant	lex	spell	resp	rt	prev_rt	rt_raw
##		<dbl>	<chr>	<chr>	<chr>	<dbl>	<dbl>
## 1	20	N	staud	N	977	511	977
## 2	9	N	dinbuss	N	565	765	565
## 3	47	N	snilling	N	562	496	562
## 4	103	N	gancens	N	572	656	572
## 5	45	W	filled	W	659	981	659
## 6	73	W	journals	W	538	1505	538
## 7	24	W	apache	W	626	546	626
## 8	11	W	flake	W	566	717	566
## 9	32	W	reliefs	W	922	1471	922
## 10	96	N	sarves	N	555	806	555

```
## # ... with 990 more rows
```

## *Renaming some with rename\_at*

In the following example, we select all variables whose names contain `rt` at their start or end, and then replace their occurrences of `rt` with `reaction_time`.

```
rename_at(bl_p_df,  
  vars(matches('^rt|rt$')),  
  ~str_replace_all(., 'rt', 'reaction_time'))
```

```
## # A tibble: 1,000 x 7
```

```
##   participant lex    spell resp  reaction_time prev.reaction_t~ reaction_time.r~  
##           <dbl> <chr> <chr> <chr>           <dbl>           <dbl>           <dbl>  
## 1           20 N      staud N              977              511              977  
## 2            9 N      dinb~ N              565              765              565  
## 3           47 N      snil~ N              562              496              562  
## 4          103 N      ganc~ N              572              656              572  
## 5           45 W      fill~ W              659              981              659  
## 6           73 W      jour~ W              538             1505              538  
## 7           24 W      apac~ W              626              546              626  
## 8           11 W      flake W              566              717              566  
## 9           32 W      reli~ W              922             1471              922  
## 10          96 N      sarv~ N              555              806              555  
## # ... with 990 more rows
```

## *Rename some with `rename_if`*

Using `rename_if`, for example, if we wanted to capitalize the names of those variables that are character variables, we could do the following.

```
rename_if(blp_df, is.character, str_to_upper)
```

```
## # A tibble: 1,000 x 7
```

```
##   participant LEX    SPELL    RESP    rt prev.rt rt.raw
##   <dbl> <chr> <chr>    <chr> <dbl>    <dbl>    <dbl>
## 1         20 N    staud    N     977      511     977
## 2          9 N    dinbuss  N     565      765     565
## 3         47 N    snilling N     562      496     562
## 4        103 N    gancens  N     572      656     572
## 5         45 W    filled   W     659      981     659
## 6         73 W    journals W     538     1505     538
## 7         24 W    apache   W     626      546     626
## 8         11 W    flake    W     566      717     566
## 9         32 W    reliefs  W     922     1471     922
## 10        96 N    sarves   N     555      806     555
```

```
## # ... with 990 more rows
```

## *Slicing data frames*

We use `slice` to select observations by their indices.

```
slice(blp_df, c(10, 20, 50, 100, 500))
```

```
## # A tibble: 5 x 7
```

```
##   participant lex    spell resp      rt prev.rt rt.raw  
##           <dbl> <chr> <chr> <chr> <dbl>   <dbl> <dbl>  
## 1           96 N      sarves N      555     806    555  
## 2           46 W      mirage W      778     571    778  
## 3           72 N      gright N      430     675    430  
## 4            3 W      gleam  W      361     370    361  
## 5           92 W      coaxes W      699     990    699
```

## Slice: Example 2

We can select consecutive ranges as follows.

```
slice(blp_df, 10:100)
```

```
## # A tibble: 91 x 7
##   participant lex    spell    resp    rt prev.rt rt.raw
##         <dbl> <chr> <chr>    <chr> <dbl>    <dbl> <dbl>
## 1         96 N    sarves    N      555      806    555
## 2         82 W    deceits    W      657      728    657
## 3         37 W    nothings    N      NA      552    712
## 4         52 N    chuespies    N      427      539    427
## 5         96 N    mowny      N     1352     1020   1352
## 6         96 N    cranned    N      907      573    907
## 7         89 N    flud      N      742      834    742
## 8          3 N    bromble    N      523      502    523
## 9          7 N    trubbles    N      782      458    782
## 10        35 N    playfound    N      643      663    643
## # ... with 81 more rows
```

## Slice: Example 3

Drop the first 10 rows.

```
slice(bl_p_df, -(1:10))
```

```
## # A tibble: 990 x 7
```

```
##   participant lex    spell    resp    rt prev.rt rt.raw
##         <dbl> <chr> <chr>    <chr> <dbl>   <dbl>  <dbl>
## 1          82 W    deceits  W      657     728    657
## 2          37 W    nothings N      NA     552    712
## 3          52 N    chuespies N     427     539    427
## 4          96 N    mowny    N    1352    1020   1352
## 5          96 N    cranned  N     907     573    907
## 6          89 N    flud     N     742     834    742
## 7           3 N    bromble  N     523     502    523
## 8           7 N    trubbles N     782     458    782
## 9          35 N    playfound N     643     663    643
## 10         46 W    mirage   W     778     571    778
## # ... with 980 more rows
```



## Slice: Example 4

We can use `n()` to indicate the last row, as in the following example.

```
slice(blp_df, 600:n())
```

```
## # A tibble: 401 x 7
##   participant lex    spell      resp      rt prev.rt rt.raw
##         <dbl> <chr> <chr>      <chr> <dbl>    <dbl> <dbl>
## 1          16 W    earthworms W        767      659    767
## 2          50 W    markers   W        664      852    664
## 3          35 N    spoton    N        522      721    522
## 4          88 W    tawny     N         NA      535    856
## 5          51 N    gember    N        562      598    562
## 6          63 W    classed   W        706      429    706
## 7          63 N    clallers  N        401      495    401
## 8           8 W    pauper    W        734     1126    734
## 9           2 W    badges    W        485      498    485
## 10         97 N    foarded   N        802      464    802
## # ... with 391 more rows
```

## Slice: Example 5

```
slice(blp_df, (n()-10):n())
```

```
## # A tibble: 11 x 7
```

##	participant	lex	spell	resp	rt	prev.rt	rt.raw
##		<dbl>	<chr>	<chr>	<chr>	<dbl>	<dbl>
## 1	29	N	khandles	N	511	777	511
## 2	88	N	ixcurs	N	504	552	504
## 3	50	N	homply	N	518	583	518
## 4	103	W	baste	W	683	454	683
## 5	67	W	tall	W	476	572	476
## 6	45	W	gardens	W	586	1023	586
## 7	105	W	goldfinch	N	NA	903	775
## 8	72	W	varmint	N	NA	507	653
## 9	3	W	lurked	W	537	520	537
## 10	3	W	village	W	538	522	538
## 11	17	W	fudge	W	410	437	410

## Filtering data frames

The `filter` command is a powerful means to filter observations according to their values. For example, we can select all the observations where the `lex` variable is `N` as follows.

```
filter(blp_df, lex == 'N')
```

```
## # A tibble: 502 x 7
```

```
##   participant lex    spell    resp    rt prev.rt rt.raw
##   <dbl> <chr> <chr>    <chr> <dbl>    <dbl> <dbl>
## 1         20 N    staud    N      977      511    977
## 2          9 N    dinbuss  N      565      765    565
## 3         47 N    snilling N      562      496    562
## 4        103 N    gancens  N      572      656    572
## 5         96 N    sarves   N      555      806    555
## 6         52 N    chuespies N      427      539    427
## 7         96 N    mowny    N     1352     1020   1352
## 8         96 N    cranned  N      907      573    907
## 9         89 N    flud     N      742      834    742
## 10        3 N    bromble  N      523      502    523
## # ... with 492 more rows
```

## Filter: Example 2

```
filter(blp_df, lex == 'N', resp=='W')
```

```
## # A tibble: 35 x 7
```

```
##   participant lex    spell    resp    rt prev.rt rt.raw
##         <dbl> <chr> <chr>    <chr> <dbl>    <dbl>  <dbl>
## 1           73 N    bunding  W      NA      978    1279
## 2           63 N    gallays  W      NA      589     923
## 3           50 N    droper   W      NA      741     573
## 4            6 N    flooder  W      NA      524     557
## 5           73 N    khantum  W      NA      623    1355
## 6           81 N    seaped   W      NA      765     691
## 7           43 N    gafers   W      NA      556     812
## 8          101 N    winchers W      NA      632     852
## 9           81 N    flaged   W      NA      674     609
## 10          11 N    frocker  W      NA      653     665
## # ... with 25 more rows
```

## Filter: Example 3

```
filter(blp_df, lex == 'N', resp=='W', rt.raw <= 500)
```

```
## # A tibble: 5 x 7
```

	participant	lex	spell	resp	rt	prev.rt	rt.raw
		<dbl>	<chr>	<chr>	<chr>	<dbl>	<dbl>
## 1	28	N	cown	W	NA	680	498
## 2	17	N	beeched	W	NA	450	469
## 3	29	N	conform	W	NA	495	497
## 4	35	N	blear	W	NA	592	461
## 5	89	N	stumming	W	NA	571	442

## Filter: Example 4

This command is equivalent to making a conjunction of conditions using & as follows.

```
filter(blp_df, lex == 'N' & resp=='W' & rt.raw <= 500)
```

```
## # A tibble: 5 x 7
```

	participant	lex	spell	resp	rt	prev.rt	rt.raw
		<dbl>	<chr>	<chr>	<chr>	<dbl>	<dbl>
## 1	28	N	cown	W	NA	680	498
## 2	17	N	beeched	W	NA	450	469
## 3	29	N	conform	W	NA	495	497
## 4	35	N	blear	W	NA	592	461
## 5	89	N	stumming	W	NA	571	442

## Filter: Example 5

We can make a *disjunction* of conditions for filtering using the logical-or symbol `|`.

```
filter(blp_df, rt.raw < 500 | rt.raw > 1000)
```

```
## # A tibble: 296 x 7
```

```
##   participant lex    spell      resp      rt prev.rt rt.raw
##   <dbl> <chr> <chr>      <chr> <dbl>   <dbl>   <dbl>
## 1         52 N    chuespies  N      427     539     427
## 2         96 N    mowny      N     1352    1020    1352
## 3         28 W    stelae     N        NA     678     497
## 4         85 W    forewarned N        NA     525     350
## 5         24 W    owl      W      470     535     470
## 6         97 W    soda       W      436     447     436
## 7         81 N    fugate     N      425     403     425
## 8        105 N    pamps      N        NA     884    1494
## 9         27 W    outgrowth  N        NA     633    1014
## 10        82 W    kitty      W      431     476     431
## # ... with 286 more rows
```

## Filter: Example 6

```
filter(blp_df, rt.raw %in% 500:510)
```

```
## # A tibble: 26 x 7
```

##	participant	lex	spell	resp	rt	prev.rt	rt.raw	
##		<dbl>	<chr>	<chr>	<chr>	<dbl>	<dbl>	
##	1	44	W	subscribed	W	509	475	509
##	2	89	W	snatcher	W	506	1004	506
##	3	2	N	tronculling	N	508	490	508
##	4	43	N	trabnate	N	510	542	510
##	5	75	N	dousleens	N	508	924	508
##	6	94	W	strangeness	W	508	522	508
##	7	68	W	greed	W	505	653	505
##	8	32	N	krifo	N	508	607	508
##	9	2	W	tweaks	W	508	474	508
##	10	85	N	waffs	N	506	471	506

```
## # ... with 16 more rows
```



## Filter: Example 7

```
filter(blp_df,  
       lex == 'W',  
       str_length(spell) < 5 & (resp != lex | rt.raw > 900))
```

```
## # A tibble: 14 x 7
```

##	participant	lex	spell	resp	rt	prev.rt	rt.raw	
##		<dbl>	<chr>	<chr>	<chr>	<dbl>	<dbl>	
##	1	21	W	bosk	N	NA	608	1532
##	2	68	W	wily	N	NA	723	636
##	3	30	W	sew	N	NA	473	524
##	4	34	W	jibs	N	NA	781	756
##	5	85	W	rote	N	NA	505	458
##	6	13	W	oofs	N	NA	560	654
##	7	72	W	awed	N	NA	1203	1801
##	8	14	W	yids	N	NA	625	620
##	9	68	W	oho	N	NA	633	630
##	10	103	W	carl	N	NA	1046	1042
##	11	46	W	brae	N	NA	644	720
##	12	81	W	bloc	N	NA	759	575
##	13	75	W	kind	W	903	1067	903

## Filter: Example 8

Using `filter_all`, we can filter rows that contain at least one NA value.

```
filter_all(bl_p_df, any_vars(is.na(.)))
```

```
## # A tibble: 179 x 7
```

```
##   participant lex    spell      resp      rt prev.rt rt.raw
##   <dbl> <chr> <chr>      <chr> <dbl>    <dbl> <dbl>
## 1      37 W    nothings    N      NA      552    712
## 2      28 W    stelae     N      NA      678    497
## 3      85 W    forewarned N      NA      525    350
## 4     105 N    pamps      N      NA     884   1494
## 5      27 W    outgrowth  N      NA     633   1014
## 6      89 W    chards     N      NA     545    754
## 7      63 N    shrudule   N      NA       0   2553
## 8      73 W    chiggers   N      NA     726    654
## 9      73 N    bunding    W      NA     978   1279
## 10     22 W    aitches    N      NA     521    665
## # ... with 169 more rows
```

## Filter: Example 9

For example, the following filters all observations where the value of all variables that start or end with `rt` are greater than 500.

```
filter_at(blp_df, vars(matches('^rt|rt$')), all_vars(. > 500))
```

```
## # A tibble: 530 x 7
```

	participant	lex	spell	resp	rt	prev.rt	rt.raw
	<dbl>	<chr>	<chr>	<chr>	<dbl>	<dbl>	<dbl>
## 1	20	N	staud	N	977	511	977
## 2	9	N	dinbuss	N	565	765	565
## 3	103	N	gancens	N	572	656	572
## 4	45	W	filled	W	659	981	659
## 5	73	W	journals	W	538	1505	538
## 6	24	W	apache	W	626	546	626
## 7	11	W	flake	W	566	717	566
## 8	32	W	reliefs	W	922	1471	922
## 9	96	N	sarves	N	555	806	555
## 10	82	W	deceits	W	657	728	657

```
## # ... with 520 more rows
```

## Filter: Example 10

Values of the `rt` variables that are lower than the medians of their variables.

```
filter_at(blp_df,  
          vars(matches('^rt|rt$')),  
          all_vars(. < median(., na.rm=T)))
```

```
## # A tibble: 251 x 7
```

	participant	lex	spell	resp	rt	prev.rt	rt.raw
	<dbl>	<chr>	<chr>	<chr>	<dbl>	<dbl>	<dbl>
## 1	47	N	snilling	N	562	496	562
## 2	52	N	chuespies	N	427	539	427
## 3	3	N	bromble	N	523	502	523
## 4	36	W	outposts	W	560	461	560
## 5	24	W	owl	W	470	535	470
## 6	97	W	soda	W	436	447	436
## 7	18	N	tessler	N	560	477	560
## 8	81	N	fugate	N	425	403	425
## 9	29	N	placker	N	542	558	542
## 10	82	W	kitty	W	431	476	431

```
## # with 241 more rows
```

## Filter: Example 11

Values of numeric variables that are lower than the medians of their variables.

```
filter_if(blp_df,  
          is.numeric,  
          all_vars(. < median(., na.rm=T)))
```

```
## # A tibble: 138 x 7
```

	participant	lex	spell	resp	rt	prev.rt	rt.raw
	<dbl>	<chr>	<chr>	<chr>	<dbl>	<dbl>	<dbl>
## 1	3	N	bromble	N	523	502	523
## 2	36	W	outposts	W	560	461	560
## 3	24	W	owl	W	470	535	470
## 4	18	N	tesslier	N	560	477	560
## 5	29	N	placker	N	542	558	542
## 6	6	N	checsons	N	491	555	491
## 7	19	N	jontage	N	413	471	413
## 8	44	W	snows	W	437	432	437
## 9	13	N	lavo	N	479	510	479
## 10	17	N	basyl	N	413	508	413

```
## # with 128 more rows
```

## *Adding or modifying variables*

- ▶ The `mutate` command is a very powerful tool in the `dplyr` toolbox.
- ▶ It allows us to create new variables and alter the values of existing ones.

## Mutate: Example 1

The following creates a new variable `acc`.

```
mutate(blp_df, acc = lex == resp)
```

```
## # A tibble: 1,000 x 8
```

##	participant	lex	spell	resp	rt	prev.rt	rt.raw	acc
##		<dbl>	<chr>	<chr>	<chr>	<dbl>	<dbl>	<lgl>
## 1	20	N	staud	N	977	511	977	TRUE
## 2	9	N	dinbuss	N	565	765	565	TRUE
## 3	47	N	snilling	N	562	496	562	TRUE
## 4	103	N	gancens	N	572	656	572	TRUE
## 5	45	W	filled	W	659	981	659	TRUE
## 6	73	W	journals	W	538	1505	538	TRUE
## 7	24	W	apache	W	626	546	626	TRUE
## 8	11	W	flake	W	566	717	566	TRUE
## 9	32	W	reliefs	W	922	1471	922	TRUE
## 10	96	N	sarves	N	555	806	555	TRUE

```
## # ... with 990 more rows
```

## Mutate: Example 2

```
mutate(bl_p_df, len = str_length(spell))
```

```
## # A tibble: 1,000 x 8
```

```
##   participant lex    spell    resp    rt prev.rt rt.raw  len
##         <dbl> <chr> <chr>    <chr> <dbl>   <dbl> <dbl> <int>
## 1           20 N      staud     N      977     511    977     5
## 2            9 N    dinbuss    N      565     765    565     7
## 3           47 N    snilling  N      562     496    562     8
## 4          103 N    gancens    N      572     656    572     7
## 5           45 W     filled    W      659     981    659     6
## 6           73 W   journals    W      538    1505    538     8
## 7           24 W     apache    W      626     546    626     6
## 8           11 W     flake     W      566     717    566     5
## 9           32 W    reliefs    W      922    1471    922     7
## 10          96 N     sarves    N      555     806    555     6
## # ... with 990 more rows
```



## Mutate: Example 3

```
mutate(bl_p_df,  
  acc = lex == resp,  
  fast = rt.raw < mean(rt.raw, na.rm=TRUE))
```

```
## # A tibble: 1,000 x 9
```

```
##   participant lex    spell    resp    rt prev.rt rt.raw acc  fast  
##      <dbl> <chr> <chr>    <chr> <dbl>    <dbl> <dbl> <lgl> <lgl>  
## 1         20 N      staud    N      977      511    977 TRUE FALSE  
## 2          9 N    dinbuss  N      565      765    565 TRUE  TRUE  
## 3         47 N    snilling N      562      496    562 TRUE  TRUE  
## 4        103 N    gancens  N      572      656    572 TRUE  TRUE  
## 5         45 W     filled  W      659      981    659 TRUE  TRUE  
## 6         73 W   journals  W      538    1505    538 TRUE  TRUE  
## 7         24 W    apache  W      626      546    626 TRUE  TRUE  
## 8         11 W     flake   W      566      717    566 TRUE  TRUE  
## 9         32 W    reliefs  W      922    1471    922 TRUE FALSE  
## 10        96 N    sarves   N      555      806    555 TRUE  TRUE  
## # ... with 990 more rows
```

## Mutate: Example 4

Change all variables to character vectors.

```
mutate_all(blpl_df, as.character)
```

```
## # A tibble: 1,000 x 7
```

##	participant	lex	spell	resp	rt	prev.rt	rt.raw
##	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
##	1 20	N	staud	N	977	511	977
##	2 9	N	dinbuss	N	565	765	565
##	3 47	N	snilling	N	562	496	562
##	4 103	N	gancens	N	572	656	572
##	5 45	W	filled	W	659	981	659
##	6 73	W	journals	W	538	1505	538
##	7 24	W	apache	W	626	546	626
##	8 11	W	flake	W	566	717	566
##	9 32	W	reliefs	W	922	1471	922
##	10 96	N	sarves	N	555	806	555
##	# ... with 990 more rows						

## Mutate: Example 5

Apply a log transform to all the rt variables.

```
mutate_at(blp_df, vars(matches('^rt|rt$')), log)
```

```
## # A tibble: 1,000 x 7
```

```
##   participant lex    spell    resp    rt prev.rt rt.raw
##         <dbl> <chr> <chr>    <chr> <dbl>    <dbl> <dbl>
## 1           20 N      staud     N      6.88     6.24  6.88
## 2            9 N    dinbuss    N      6.34     6.64  6.34
## 3           47 N    snilling  N      6.33     6.21  6.33
## 4          103 N    gancens    N      6.35     6.49  6.35
## 5           45 W    filled     W      6.49     6.89  6.49
## 6           73 W   journals    W      6.29     7.32  6.29
## 7           24 W    apache     W      6.44     6.30  6.44
## 8           11 W    flake      W      6.34     6.58  6.34
## 9           32 W   reliefs    W      6.83     7.29  6.83
## 10          96 N    sarves     N      6.32     6.69  6.32
## # ... with 990 more rows
```

## Mutate: Example 6

Change all character vectors to factors.

```
mutate_if(bl_p_df, is.character, as.factor)
```

```
## # A tibble: 1,000 x 7
```

```
##   participant lex    spell    resp    rt prev.rt rt.raw
##         <dbl> <fct> <fct>    <fct> <dbl>    <dbl> <dbl>
## 1           20 N      staud     N      977      511    977
## 2            9 N    dinbuss    N      565      765    565
## 3           47 N    snilling    N      562      496    562
## 4          103 N    gancens    N      572      656    572
## 5           45 W     filled    W      659      981    659
## 6           73 W   journals    W      538     1505    538
## 7           24 W    apache    W      626      546    626
## 8           11 W     flake    W      566      717    566
## 9           32 W   reliefs    W      922     1471    922
## 10          96 N     sarves    N      555      806    555
## # ... with 990 more rows
```

## Mutate: Example 8

Create a new variable `speed` that takes the value of `fast` if `rt.raw` is less than 750, and `slow` otherwise.

```
mutate(blp_df,  
      speed = if_else(rt.raw < 750,  
                      'fast',  
                      'slow')  
)
```

```
## # A tibble: 1,000 x 8
```

```
##   participant lex    spell    resp      rt prev.rt rt.raw speed  
##           <dbl> <chr> <chr>    <chr> <dbl>    <dbl> <dbl> <chr>  
## 1           20 N      staud     N      977      511    977 slow  
## 2            9 N    dinbuss    N      565      765    565 fast  
## 3           47 N    snilling  N      562      496    562 fast  
## 4          103 N    gancens    N      572      656    572 fast  
## 5           45 W     filled    W      659      981    659 fast  
## 6           73 W   journals    W      538    1505    538 fast  
## 7           24 W    apache    W      626      546    626 fast  
## 8           11 W     flake    W      566      717    566 fast  
## 9           32 W   reliefs    W      922    1471    922 slow  
## 10          96 N    sarves    N      555      806    555 fast  
## # ... with 990 more rows
```

## Mutate: Example 9

Replace the `lex` variable's values W and N with word and nonword.

```
mutate(blp_df,  
      lex = recode(lex, 'W'='word', 'N'='nonword')  
)
```

```
## # A tibble: 1,000 x 7  
##   participant lex      spell      resp      rt prev.rt rt.raw  
##         <dbl> <chr>    <chr>    <chr> <dbl>    <dbl> <dbl>  
## 1          20 nonword  staud      N      977      511    977  
## 2           9 nonword  dinbuss    N      565      765    565  
## 3          47 nonword  snilling  N      562      496    562  
## 4         103 nonword  gancens    N      572      656    572  
## 5          45 word      filled     W      659      981    659  
## 6          73 word      journals   W      538     1505    538  
## 7          24 word      apache     W      626      546    626  
## 8          11 word      flake      W      566      717    566  
## 9          32 word      reliefs    W      922     1471    922  
## 10         96 nonword  sarves     N      555      806    555  
## # ... with 990 more rows
```

## Mutate: Example 10

We could use `case_when` to convert values of `prev.rt` that are below 500 to `fast`, and those above 1500 to `slow`, and those in between 500 and 1500 to `medium`.

```
mutate(bl_p_df,  
  prev.rt = case_when(  
    prev.rt < 500 ~ 'fast',  
    prev.rt > 1500 ~ 'slow',  
    TRUE ~ 'medium'  
  )  
)
```

```
## # A tibble: 1,000 x 7
```

	participant	lex	spell	resp	rt	prev.rt	rt.raw
		<dbl>	<chr>	<chr>	<chr>	<dbl>	<chr>
##	1	20	N	staud	N	977	medium
##	2	9	N	dinbuss	N	565	medium
##	3	47	N	snilling	N	562	fast
##	4	103	N	gancens	N	572	medium
##	5	45	W	filled	W	659	medium
##	6	73	W	journals	W	538	slow
##	7	24	W	apache	W	626	medium
##	8	11	W	flake	W	566	medium
##	9	32	W	reliefs	W	922	medium
##	10	96	N	sarves	N	555	medium

```
## # ... with 990 more rows
```

## Mutate: Example 11

Use `plyr::mapvalues` to map a range of values to another.

```
mutate(blp_df,  
      rt_reverse = plyr::mapvalues(rt, from=500:1000, to=1000:500)  
)
```

```
## # A tibble: 1,000 x 8
```

```
##   participant lex    spell    resp    rt prev.rt rt.raw rt_reverse  
##           <dbl> <chr> <chr>    <chr> <dbl>    <dbl>  <dbl>    <dbl>  
## 1           20 N      staud    N      977      511    977      523  
## 2            9 N      dinbuss  N      565      765    565      935  
## 3           47 N      snilling N      562      496    562      938  
## 4          103 N      gancens  N      572      656    572      928  
## 5           45 W      filled  W      659      981    659      841  
## 6           73 W      journals W      538     1505    538      962  
## 7           24 W      apache  W      626      546    626      874  
## 8           11 W      flake   W      566      717    566      934  
## 9           32 W      reliefs W      922     1471    922      578  
## 10          96 N      sarves  N      555      806    555      945  
## # ... with 990 more rows
```



## Sorting data frames

We can sort data frames with `arrange`. For example to sort by participant and then by spell, we would do the following.

```
arrange(blp_df, participant, spell)
```

```
## # A tibble: 1,000 x 7
```

	participant	lex	spell	resp	rt	prev.rt	rt.raw
	<dbl>	<chr>	<chr>	<chr>	<dbl>	<dbl>	<dbl>
## 1	1	W	abyss	W	629	683	629
## 2	1	N	baisees	N	524	574	524
## 3	1	W	carport	W	779	605	779
## 4	1	N	cellies	N	792	652	792
## 5	1	W	chafing	W	601	720	601
## 6	1	N	dametails	N	694	635	694
## 7	1	N	foother	N	789	566	789
## 8	1	W	gantries	W	644	581	644
## 9	1	N	hogtush	N	679	568	679
## 10	1	N	lisedess	N	679	619	679

```
## # ... with 990 more rows
```

## Sorting in descending order

We can use `desc` to sort in descending order.

```
arrange(blp_df, participant, desc(spell))
```

```
## # A tibble: 1,000 x 7
```

```
##   participant lex    spell    resp    rt prev.rt rt.raw
##         <dbl> <chr> <chr>    <chr> <dbl>   <dbl>  <dbl>
## 1           1 N    wintes    N     545     629    545
## 2           1 N    treeps    N     607     610    607
## 3           1 W    squashes W     494     491    494
## 4           1 N    sinkhicks N     536     519    536
## 5           1 W    shafting W     553     571    553
## 6           1 W    month    W     500     498    500
## 7           1 N    lisedess N     679     619    679
## 8           1 N    hogtush  N     679     568    679
## 9           1 W    gantries W     644     581    644
## 10          1 N    footer  N     789     566    789
## # ... with 990 more rows
```

## Summarizing variables

We can use `summarize` to calculate some summary statistics of particular variables.

```
summarize(blp_df,  
          mean_rt = mean(rt, na.rm = T),  
          median_rt = median(rt, na.rm = T),  
          sd_rt.raw = sd(rt.raw, na.rm = T)  
)
```

```
## # A tibble: 1 x 3  
##   mean_rt median_rt sd_rt.raw  
##   <dbl>     <dbl>    <dbl>  
## 1    638.       588     474.
```

## Summarize: Example 2

We can use the `summarize_all` variant of `summarize` to apply a summarisation function to all variables, as in the following example.

```
summarize_all(blp_df, n_distinct)
```

```
## # A tibble: 1 x 7
##   participant    lex spell  resp    rt prev.rt rt.raw
##         <int> <int> <int> <int> <int>   <int>  <int>
## 1           78     2  990     2  421    493   516
```

## Summarize: Example 3

In the following example, we calculate the mean of all the reaction times variables.

```
summarize_at(blp_df, vars(matches('^rt|rt$')), ~mean(., na.rm=T))
```

```
## # A tibble: 1 x 3
##       rt prev.rt rt.raw
##   <dbl>   <dbl> <dbl>
## 1  638.    660.   708.
```

## Summarize: Example 4

The `summarize_if` will apply the summary function to numeric variables.

```
summarize_if(blp_df, is.numeric, ~mean(., na.rm=T))
```

```
## # A tibble: 1 x 4  
##   participant      rt prev.rt rt.raw  
##       <dbl> <dbl>   <dbl>  <dbl>  
## 1         49.5  638.    660.   708.
```

## Summarize: Example 5

In the following, we calculate the same three summary statistics for two variables.

```
summarise_at(bl_p_df,  
  vars(rt, rt.raw),  
  list(mean = ~mean(., na.rm=T),  
        median = ~median(., na.rm=T),  
        sd = ~sd(., na.rm=T)  
  )  
)
```

```
## # A tibble: 1 x 6  
##   rt_mean rt.raw_mean rt_median rt.raw_median rt_sd rt.raw_sd  
##   <dbl>      <dbl>      <dbl>          <dbl> <dbl>      <dbl>  
## 1    638.        708.        588            605  191.      474.
```

## *The %>% pipe*

- ▶ To understand pipes, let us begin with a very simple example.
- ▶ The following is a vector of the first 10 prime numbers.

```
primes <- c(2, 3, 5, 7, 11, 13, 17, 19, 23, 29)
```

- ▶ We can calculate the sum of primes as follows.

```
sum(primes)
```

```
## [1] 129
```

- ▶ We may then calculate the square root of this sum.

```
sqrt(sum(primes))
```

```
## [1] 11.35782
```

- ▶ We may then calculate the logarithm of this square root.

```
log(sqrt(sum(primes)))
```

```
## [1] 2.429906
```



## The %>% pipe

- Using the pipe %>%, we can replace

```
log(sqrt(sum(primes)))
```

```
## [1] 2.429906
```

with

```
primes %>% sum() %>% sqrt() %>% log()
```

```
## [1] 2.429906
```

- In general, if we have a variable  $x$  and a function  $f()$ , then  $f(x)$  is equivalent to the following.

```
x %>% f()    # equivalent to f(x)
```

## Grouped summaries

- ▶ Using `group_by` and `summarize` together is a powerful way to create new (reduced) data frames.
- ▶ For example, for each participant, and for each of the two stimuli types (i.e. the N and W values of `lex`), we can calculate summary statistics.

```
group_by(blp_df, participant, lex) %>%
summarize(n_stimuli = n(),
          correct_resp = sum(resp == lex, na.rm=T),
          reaction_time = mean(rt.raw, na.rm=T))
```

  

```
## # A tibble: 156 x 5
## # Groups:   participant [78]
##   participant lex   n_stimuli correct_resp reaction_time
##         <dbl> <chr>      <int>         <int>         <dbl>
## 1             1 N             9             9           649.
## 2             1 W             7             7           600
## 3             2 N             7             6           625.
## 4             2 W             6             5           477.
## 5             3 N             4             4           540.
## 6             3 W             8             7           529
## 7             4 N             5             5           589.
## 8             4 W             5             4           465.
## 9             5 N             1             1           495
## 10            5 W             3             2           571
## # ... with 146 more rows
```

## Pipelines: Example 1

```
blp_df %>%  
  mutate(accuracy = resp == lex,  
         stimulus = recode(lex, 'W'='word', 'N'='nonword')  
  ) %>%  
  select(participant, stimulus, item=spell, accuracy, speed=rt.raw) %>%  
  arrange(participant, speed)
```

```
## # A tibble: 1,000 x 5  
##   participant stimulus item      accuracy speed  
##   <dbl> <chr>      <chr>      <lgl>      <dbl>  
## 1           1 word      squashes  TRUE        494  
## 2           1 word      month     TRUE        500  
## 3           1 nonword  baiseses  TRUE        524  
## 4           1 nonword  sinkhicks TRUE        536  
## 5           1 nonword  wintes    TRUE        545  
## 6           1 word      shafting  TRUE        553  
## 7           1 word      chafing   TRUE        601  
## 8           1 nonword  treeps    TRUE        607  
## 9           1 word      abyss     TRUE        629  
## 10          1 word      gantries  TRUE        644  
## # ... with 990 more rows
```

## Pipelines: Example 2

```
blp_df %>%  
  filter(lex == 'W') %>%  
  mutate(word_length = str_length(spell),  
         accuracy = resp == lex) %>%  
  rename(speed = rt.raw) %>%  
  group_by(word_length) %>%  
  summarize_at(vars(accuracy, speed), ~mean(., na.rm=T)) %>%  
  ungroup() %>%  
  select(word_length, accuracy, speed) %>%  
  arrange(word_length, accuracy, speed)
```

```
## # A tibble: 9 x 3  
##   word_length accuracy speed  
##       <int>     <dbl> <dbl>  
## 1         3      0.7    551.  
## 2         4     0.744   649.  
## 3         5     0.718   825.  
## 4         6     0.807   723.  
## 5         7     0.821   704.  
## 6         8     0.835   678.  
## 7         9     0.595   914.  
## 8        10     0.714   670.  
## 9        11     0.5    700.
```

## Combining data frames

- ▶ There are two main ways to combine data frames: *binds* and *joins*.
- ▶ A *bind* operation is a simple operation that either vertically stack data frames that share common variables, or horizontally stack data frames that have the same number of observations.
- ▶ A *join* merges data frames by shared variables.

## *Binds*

- To illustrate, we will create three small data frames.

```
Df_1 <- tibble(x = c(1, 2, 3),  
               y = c(2, 7, 1),  
               z = c(0, 2, 7))
```

```
Df_2 <- tibble(y = c(5, 7),  
               z = c(6, 7),  
               x = c(1, 2))
```

```
Df_3 <- tibble(a = c(5, 6, 1),  
               b = c('a', 'b', 'c'),  
               c = c(T, T, F))
```

## *Bind rows*

- ▶ The Df\_1 and Df\_2 data frames share common variable names.
- ▶ They can be vertically stacked using a bind\_rows operation.

```
bind_rows(Df_1, Df_2)
```

```
## # A tibble: 5 x 3
##       x     y     z
##   <dbl> <dbl> <dbl>
## 1     1     2     0
## 2     2     7     2
## 3     3     1     7
## 4     1     5     6
## 5     2     7     7
```

## Bind columns

- The Df\_1 and Df\_3 data frames have the same number of observations and so can be stacked side by side with a `bind_cols` operation.

```
bind_cols(Df_1, Df_3)
```

```
## # A tibble: 3 x 6
##       x     y     z     a b     c
##   <dbl> <dbl> <dbl> <dbl> <chr> <lgl>
## 1     1     2     0     5 a    TRUE
## 2     2     7     2     6 b    TRUE
## 3     3     1     7     1 c    FALSE
```



## Combining data frames by joins

- ▶ A *join* operation is a common operation in relational databases using SQL.
- ▶ It allows us to join separate tables according to shared keys.
- ▶ As an example, consider `stimuli`. It has a `spell` variable like `blp_df` file, and an three additional variables.

```
stimuli <- read_csv('data/blp_stimuli.csv')
stimuli
```

```
## # A tibble: 55,865 x 4
##   spell    old20    bnc subtex
##   <chr>    <dbl> <dbl>   <dbl>
## 1 a/c      1.95    14      0
## 2 aas      1.55     9      1
## 3 aback    1.85   327     15
## 4 abaft     2      8       2
## 5 aband    1.95     0      0
## 6 abase     1.7     6       2
## 7 abased    1.75     6      0
## 8 abashed   1.85    57      0
## 9 abate     1.75    69      5
## 10 abates   1.75     9      2
## # ... with 55,855 more rows
```

## *Join with inner\_join*

- ▶ We can join `blp_df` and `stimuli` with `inner_join`.
- ▶ An `inner_join` operation, searches through the values of variables that are shared by the two data frames in order to find matching values.
- ▶ In `blp_df` and `stimuli`, there is just one shared variable, namely `spell`.
- ▶ Thus, an `inner_join` of `blp_df` and `stimuli` will find values of `spell` on the left hand data frame that occur as values of `spell` on the right hand side.

## Join with `inner_join`

```
inner_join(blp_df, stimuli)
```

```
## # A tibble: 1,000 x 10
```

```
##   participant lex    spell    resp    rt prev.rt rt.raw old20    bnc
##   <dbl> <chr> <chr>    <chr> <dbl>    <dbl> <dbl> <dbl> <dbl>
## 1         20 N      staud    N      977      511    977  1.85     0
## 2          9 N    dinbuss  N      565      765    565  2.9      0
## 3         47 N    snilling N      562      496    562  1.8      0
## 4        103 N    gancens  N      572      656    572  2.3      0
## 5         45 W    filled  W      659      981    659  1.45  5340
## 6         73 W   journals W      538    1505    538  2.7   1030
## 7         24 W    apache  W      626      546    626  2.45   130
## 8         11 W    flake   W      566      717    566  1.5    274
## 9         32 W   reliefs W      922    1471    922  2.25   185
## 10        96 N    sarves  N      555      806    555  1.65     0
## # ... with 990 more rows
```

## *Inner, Left, right, full joins*

- ▶ In general, in an `inner_join`, if the left hand data frame has no values on the shared variables that match those on the right hand data frame, the observations from the left hand data frame are dropped.
- ▶ In addition, all observations on the right hand data frame that do not have matching observations on the left always get dropped too. However, consider the following two data frames.

```
Df_a <- tibble(x = c(1, 2, 3),  
               y = c('a', 'b', 'c'))  
Df_b <- tibble(x = c(2, 3, 4),  
               z = c('d', 'e', 'f'))
```

## *Inner, Left, right, full joins*

- The first value of x in Df\_a does not match any value of x in Df\_b, and so the corresponding observation is dropped in an inner\_join.

```
inner_join(Df_a, Df_b)
```

```
## # A tibble: 2 x 3
##       x y      z
##   <dbl> <chr> <chr>
## 1     2 b      d
## 2     3 c      e
```

## Left join

- ▶ A `left_join`, on the other hand, will preserve all values on the left and put NA as the corresponding values of the right's variables if there are no matching values.

```
left_join(Df_a, Df_b)
```

```
## # A tibble: 3 x 3
##       x y      z
##   <dbl> <chr> <chr>
## 1     1  a    <NA>
## 2     2  b      d
## 3     3  c      e
```

## *Right join*

- ▶ A `right_join` preserves all observations from the right, and places NA as the corresponding values of variables from the left that are not matched.

```
right_join(Df_a, Df_b)
```

```
## # A tibble: 3 x 3  
##       x y      z  
##   <dbl> <chr> <chr>  
## 1     2 b      d  
## 2     3 c      e  
## 3     4 <NA>    f
```

## Full joins

- A `full_join` keeps all observation in both the left and right data frames.

```
full_join(Df_a, Df_b)
```

```
## # A tibble: 4 x 3
##       x y      z
##   <dbl> <chr> <chr>
## 1     1  1 a    <NA>
## 2     2  2 b      d
## 3     3  3 c      e
## 4     4  4 <NA>   f
```



## *Specifying the by variables*

- Consider the following cases.

```
Df_4 <- tibble(x = c(1, 2, 3),  
              y = c(2, 7, 1),  
              z = c(0, 2, 7))
```

```
Df_5 <- tibble(a = c(1, 1, 7),  
              b = c(2, 3, 7),  
              c = c('a', 'b', 'c'))
```

## *Specifying the by variables*

- In the following example, we look for matches between x and y on the left and a and b on the right.

```
inner_join(Df_4, Df_5, by=c('x' = 'a', 'y' = 'b'))
```

```
## # A tibble: 1 x 4
##       x       y       z c
##   <dbl> <dbl> <dbl> <chr>
## 1     1     2     0 a
```

## Reshaping with *pivot\_longer* and *pivot\_wider*

- ▶ A so-called *tidy* data set is a data set where all rows are observations, all columns are variables, and each variable is a single value.
- ▶ Consider the following data frame.

```
recall_df <- read_csv('data/repeated_measured_a.csv')  
recall_df
```

```
## # A tibble: 5 x 4  
##   Subject    Neg    Neu    Pos  
##   <chr>    <dbl> <dbl> <dbl>  
## 1 Faye      26     12    42  
## 2 Jason     29      8    35  
## 3 Jim       32     15    45  
## 4 Ron       22     10    38  
## 5 Victor    30     13    40
```

- ▶ In this data frame, the Neg, Neu, Pos are, in fact, *values* of a variable, namely the condition of the experiment.

## *pivot\_longer*

```
recall_long <- pivot_longer(recall_df,  
                             cols = -Subject,  
                             names_to = 'condition',  
                             values_to = 'score')  
  
recall_long
```

```
## # A tibble: 15 x 3  
##   Subject condition score  
##   <chr>    <chr>    <dbl>  
## 1 Faye     Neg         26  
## 2 Faye     Neu         12  
## 3 Faye     Pos         42  
## 4 Jason    Neg         29  
## 5 Jason    Neu          8  
## 6 Jason    Pos         35  
## 7 Jim      Neg         32  
## 8 Jim      Neu         15  
## 9 Jim      Pos         45  
## 10 Ron     Neg         22  
## 11 Ron     Neu         10  
## 12 Ron     Pos         38  
## 13 Victor  Neg         30  
## 14 Victor  Neu         13  
## 15 Victor  Pos         40
```

## *pivot\_wider*

- The inverse of a `pivot_longer` is a `pivot_wider`.

```
pivot_wider(recall_long, names_from = 'condition', values_from =
```

```
## # A tibble: 5 x 4
##   Subject    Neg    Neu    Pos
##   <chr>    <dbl> <dbl> <dbl>
## 1 Faye      26     12    42
## 2 Jason     29      8    35
## 3 Jim       32     15    45
## 4 Ron       22     10    38
## 5 Victor    30     13    40
```

## *pivot\_longer: example 2*

Consider the following data.

```
recall_2_df <- read_csv('data/repeated_measured_b.csv')  
recall_2_df
```

```
## # A tibble: 5 x 7
```

```
##   Subject Cued_Neg Cued_Neu Cued_Pos Free_Neg Free_Neu Free_P  
##   <chr>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <db  
## 1 Faye          15         16         14         13         13  
## 2 Jason          4          9         10          6          7  
## 3 Jim           7          9         10          8          9  
## 4 Ron          17         18         20         12         14  
## 5 Victor        16         13         14         12         13
```

## *pivot\_longer: example 2*

```
pivot_longer(recall_2_df,  
             cols = -Subject,  
             names_to = c('cue', 'emotion'),  
             names_pattern = '(Cued|Free)_(Neg|Pos|Neu)',  
             values_to = 'score')
```

```
## # A tibble: 30 x 4  
##   Subject cue    emotion score  
##   <chr>   <chr> <chr>    <dbl>  
## 1 Faye    Cued   Neg      15  
## 2 Faye    Cued   Neu      16  
## 3 Faye    Cued   Pos      14  
## 4 Faye    Free   Neg      13  
## 5 Faye    Free   Neu      13  
## 6 Faye    Free   Pos      12  
## 7 Jason   Cued   Neg       4  
## 8 Jason   Cued   Neu       9  
## 9 Jason   Cued   Pos      10  
## 10 Jason  Free   Neg       6  
## # ... with 20 more rows
```

## *pivot\_longer: example 3*

- Consider the following data frame from earlier.

```
tmp_df <- summarise_at(blp_df,  
  vars(rt:rt.raw),  
  list(avg = ~mean(., na.rm = T),  
        med = ~median(., na.rm = T),  
        stdev = ~sd(., na.rm = T))  
)  
tmp_df
```

```
## # A tibble: 1 x 9  
##   rt_avg prev.rt_avg rt.raw_avg rt_med prev.rt_med rt.raw_med rt_stdev  
##   <dbl>      <dbl>      <dbl>  <dbl>      <dbl>      <dbl>    <dbl>  
## 1   638.        660.        708.   588        594        605    191.  
## # ... with 2 more variables: prev.rt_stdev <dbl>, rt.raw_stdev <dbl>
```



## *pivot\_longer: example 3*

```
tmp_df %>% pivot_longer(cols = everything(),  
                        names_to = c('variable', '.value'),  
                        names_pattern = "(.*)_(.*)")
```

```
## # A tibble: 3 x 4  
##   variable    avg    med stdev  
##   <chr>    <dbl> <dbl> <dbl>  
## 1 rt          638.   588  191.  
## 2 prev.rt     660.   594  253.  
## 3 rt.raw      708.   605  474.
```